

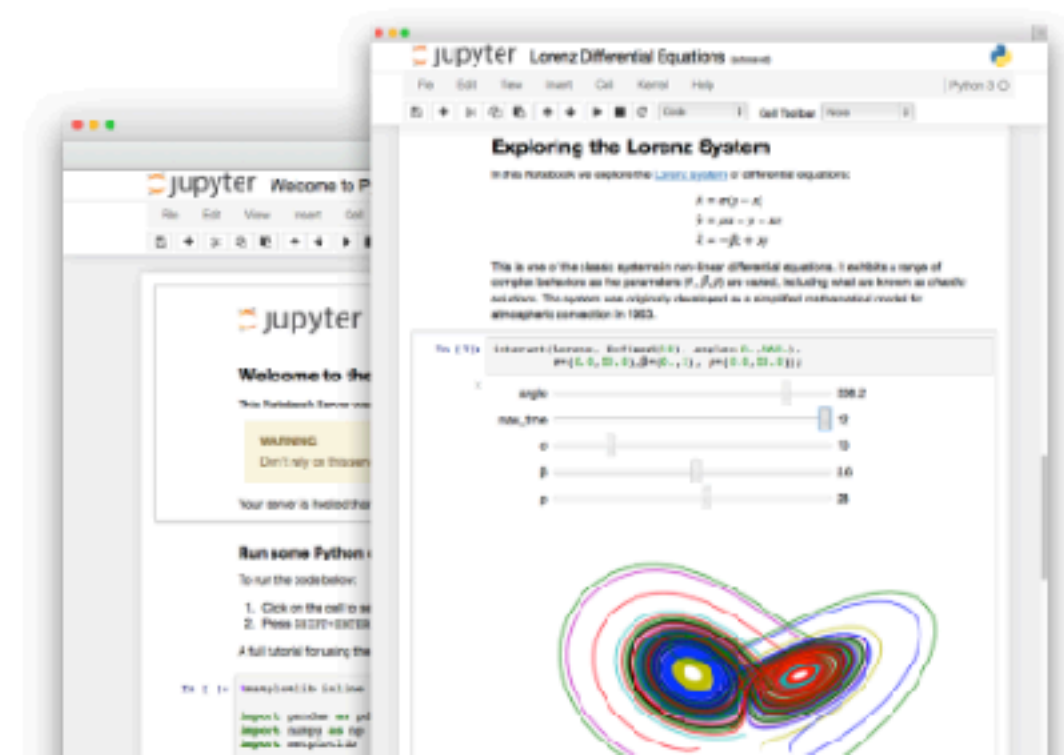
Linear Regression with Python

Tools



Jupyter Notebook

The Jupyter Notebook is a web-based interactive computing platform that allows users to author data- and code-driven narratives that combine live code, equations, narrative text, visualizations, interactive dashboards and other media.



Projects/jupyter-notebook/

Sebastian

localhost:8888/tree/Projects/jupyter-notebook#notebooks

jupyter

Logout

Files

Running

Clusters

Select items to perform actions on them.

Upload

New

☐

/ Projects / jupyter-notebook

Name

Last Modified

..

seconds ago

☐

my-first-notebook.ipynb

6 days ago

Projects/jupyter-notebook/Untitled

localhost:8888/notebooks/Projects/jupyter-notebook/Untitled.ipynb?kernel_name=python3

jupyter

Untitled

Last Checkpoint: a few seconds ago (unsaved changes)

Logout

FileEditViewInsertCellKernelWidgetsHelpTrustedPython 3

In []:

```
In [1]: print('Hello World')
```

```
Hello World
```


```
In [2]: i = 1  
while i <= 10:  
    print(i)  
    i = i + 1
```

```
1  
2  
3  
4  
5  
6  
7  
8  
9  
10
```

```
In [ ]:
```

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



pandas

- open source library
- high-performance, easy-to-use data structures and data analysis tools

<https://pandas.pydata.org/>



NumPy

[Scipy.org](#)

NumPy

NumPy is the fundamental package for scientific computing with Python. It contains among other things:

- a powerful N-dimensional array object
- sophisticated (broadcasting) functions
- tools for integrating C/C++ and Fortran code
- useful linear algebra, Fourier transform, and random number capabilities

Besides its obvious scientific uses, NumPy can also be used as an efficient multi-dimensional container of generic data. Arbitrary data-types can be defined. This allows NumPy to seamlessly and speedily integrate with a wide variety of databases.

NumPy is licensed under the [BSD license](#), enabling reuse with few restrictions.

Getting Started

- [Getting NumPy](#)
- [Installing the SciPy Stack](#)
- [NumPy and SciPy documentation page](#)
- [NumPy Tutorial](#)
- [NumPy for MATLAB® Users](#)
- [NumPy functions by category](#)
- [NumPy Mailing List](#)

For more information on the SciPy Stack (for which NumPy provides the fundamental array data structure), see [scipy.org](#).

[Donate to NumPy](#)

[Open Hub](#)

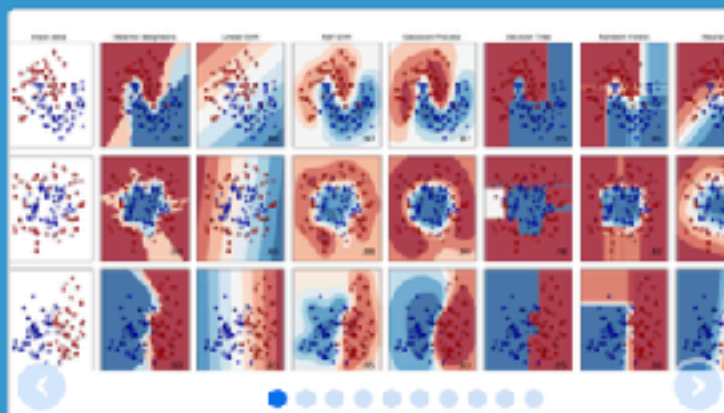
[NumPy](#)

[About NumPy](#)

[License](#)

[Old array packages](#)

<http://www.numpy.org/>



scikit-learn

Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ... — [Examples](#)

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, ridge regression, Lasso, ... — [Examples](#)

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, ... — [Examples](#)

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization. — [Examples](#)

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning

Modules: grid search, cross validation, metrics. — [Examples](#)

Preprocessing

Feature extraction and normalization.


Application: Transforming input data such as text for use with machine learning algorithms.

Modules: preprocessing, feature extraction. — [Examples](#)


<http://scikit-learn.org/stable/>

Linear Regression

$$y = a x + b$$



$$y = a x + b y + c$$



General formulation:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_D x_D$$


General formulation:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_D x_D$$

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j x_j$$

General formulation:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_D x_D$$



?

The key property of this model is that it is a linear function of the parameters !

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j x_j \quad \text{Linear in } \mathbf{x} \text{ and } \mathbf{w}$$



$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) \quad \text{Linear in } \mathbf{w} \text{ and } \phi(\mathbf{x})$$

The key property of this model is that it is a linear function of the parameters !

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j x_j \quad \text{Linear in } \mathbf{x} \text{ and } \mathbf{w}$$

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) \quad \text{Linear in } \mathbf{w} \text{ and } \phi(\mathbf{x})$$

Basis functions



$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

$$\phi_0(\mathbf{x}) = 1$$

$$y(\mathbf{x}, \mathbf{w}) = \mathbf{w}^\top \boldsymbol{\phi}(\mathbf{x})$$

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j x_j$$

- **goal:** Predicting a continuous y as a function of the variables \mathbf{x} and the parameters \mathbf{w}
- **Model assumption:**
 - data generated by $t = y(\mathbf{x}, \mathbf{w}) + \epsilon$
 - Error $\epsilon \sim$ normal distribution with zero mean
- **Loss function:** sum-of-squares error between t and prediction y (least squares)